ricu: **R**'s Interface to Intensive Care Data

Nicolas Bennett[*]
ETH Zürich

Drago Plečko[*]
ETH Zürich

Ida-Fong Uko
Monash Health

Nicolai Meinshausen
ETH Zürich

Peter Bühlmann
ETH Zürich

Abstract

Providing computational infrastructure for handling diverse intensive care unit (ICU) datasets, the R package **ricu** enables writing dataset-agnostic analysis code, thereby facilitating multi-center training and validation of machine learning models. The package is designed with an emphasis on extensibility both to new datasets as well as clinical data concepts, and currently supports the loading of around 100 patient variables corresponding to a total of 319,402 ICU admissions from 4 data sources collected in Europe and the United States. By allowing for the addition of user-specified medical concepts and data sources the aim of **ricu** is to foster robust, data-based intensive care research, allowing the user to externally validate their method or conclusion with relative ease, and in turn facilitating reproducible and therefore transparent work in this field.

Keywords: electronic health records, computational physiology, critical care medicine.

```
Note: Examples in this vignette require that one or more of datasets '
mimic ', ' mimic_demo ', ' eicu ', ' eicu_demo ', ' hirid ', ' aumc '
are available. Chunks that depend on certain datasets will not be
evaluated if the corresponding dataset is missing. In order to download
and setup data, have a look at '?setup_src_data'.
```

1. Introduction

Collection of electronic health records has seen a significant rise in recent years (Evans 2016), opening up opportunities and providing the grounds for a large body of data-driven research oriented towards helping clinicians in decision-making and therefore improving patient care and health outcomes (Jiang, Jiang, Zhi, Dong, Li, Ma, Wang, Dong, Shen, and Wang 2017).

---

[*]These authors contributed equally.

One example of a problem that has received much attention from the machine learning community is early prediction of sepsis in ICU (Desautels, Calvert, Hoffman, Jay, Kerem, Shieh, Shimabukuro, Chettipally, Feldman, Barton *et al.* 2016; Nemati, Holder, Razmi, Stanley, Clifford, and Buchman 2018; Futoma, Hariharan, Sendak, Brajer, Clement, Bedoya, O'Brien, and Heller 2017; Kam and Kim 2017). Interestingly, there is evidence that a large proportion of the publications are based on the same dataset (Fleuren, Klausch, Zwager, Schoonmade, Guo, Roggeveen, Swart, Girbes, Thoral, Ercole, Hoogendoorn, and Elbers 2020), the Medical Information Mart for Intensive Care III (MIMIC-III; Johnson, Pollard, Shen, Li-wei, Feng, Ghassemi, Moody, Szolovits, Celi, and Mark 2016), which shows a systematic lack of external validation. Part of this problem might well be the need for computational infrastructure handling multiple datasets. The MIMIC-III dataset consists of 26 different tables containing about 20GB of data. While much work and care has gone into data preprocessing in order to provide a self-contained ready-to-use data resource with MIMIC-III, seemingly simple tasks such as computing a sepsis-related organ failure assessment (SOFA) score (Vincent, Moreno, Takala, Willatts, De Mendonça, Bruining, Reinhart, Suter, and Thijs 1996) remains a non-trivial effort[1]. This is only exacerbated when aiming to co-integrate multiple different datasets of this form, spanning hospitals and even countries, in order to capture effects of differing practice and demographics.

The aim of the **ricu** package is to provide computational infrastructure allowing users to investigate complex research questions in the context of critical care medicine as easily as possible by providing a unified interface to a heterogeneous set of data sources. The package enables users to write dataset-agnostic code which can simplify implementation and shorten the time necessary for prototyping code querying different datasets. In its current form, the package handles four large-scale, publicly available intensive care databases out of the box: MIMIC-III from the Beth Israel Deaconess Medical Center in Boston, Massachusetts (Johnson *et al.* 2016), the eICU Collaborative Research Database (Pollard, Johnson, Raffa, Celi, Mark, and Badawi 2018), containing data collected from 208 hospitals across the United States, the High Time Resolution ICU Dataset (HiRID) from the Department of Intensive Care Medicine of the Bern University Hospital, Switzerland (Faltys, Zimmermann, Lyu, Hüser, Hyland, Rätsch, and Merz 2021) and AmsterdamUMCdb from the Amsterdam University Medical Center (Thoral, Peppink, Driessen, Sijbrands, Kompanje, Kaplan, Bailey, Kesecioglu, Cecconi, Churpek, Clermont, van der Schaar, Ercole, Girbes, Elbers, Force, and the SCCM/ESICM Joint Data Science Task 2021). Furthermore, **ricu** was designed with extensibility in mind such that adding further public and/or private user-provided datasets is possible. Being implemented in R, a programming language popular among statisticians and data analysts, it is our hope to contribute to accessible and reproducible research by using a familiar environment and requiring only few system dependencies, thereby simplifying setup considerably.

To our knowledge, infrastructure that provides a common interface to multiple such datasets

---

[1]There is considerable heterogeneity in number of patients satisfying the Sepsis-3 criterion (of which SOFA provides a major component; Singer, Deutschman, Seymour, Shankar-Hari, Annane, Bauer, Bellomo, Bernard, Chiche, Coopersmith, Hotchkiss, Levy, Marshall, Martin, Opal, Rubenfeld, van der Poll, Vincent, and Angus 2016) among studies investigating MIMIC-III. Reported Sepsis-3 prevalence ranges from 11.3% (Desautels *et al.* 2016), over 23.9% (Nemati *et al.* 2018) and 25.4% (Wang, Sun, Schroeder, Ameko, Moore, and Barnes 2018), up to 49.1% (Johnson, Aboab, Raffa, Pollard, Deliberato, Celi, and Stone 2018). While some of this variation may be explained by differing patient inclusion criteria, diversity in label implementation must also contribute significantly.

is a novel contribution. While there have been efforts (Adibuzzaman, Musselman, Johnson, Brown, Pitluk, and Grama 2016; Wang, McDermott, Chauhan, Ghassemi, Hughes, and Naumann 2020) attempting to abstract away some specifics of a dataset, these have so far exclusively focused on MIMIC-III, the most popular of public ICU datsets and have not been designed with dataset interoperability in mind.

Given the somewhat narrow focus of the targeted datasets, combined with the fact that in some cases data is even extracted from identical patient care management systems, it may come as a surprise as to how heterogeneous the resulting datasets are. In MIMIC-III and HiRID, for example, time-stamps are reported as absolute times (albeit randomly shifted due to data privacy concerns), whereas eICU and AUMC use relative times (with origins being admission times). Another example, involves different types of patient identifiers and their use among datasets. Common to all is the notion of an ICU admission ID, but apart from that, the amount of available information varies: While ICU (and hospital) readmissions for a given patient can be identified in some, this is not possible in other datasets. Furthermore, use of identifier systems might not be consistent over tables. In MIMIC-III, for example, some tables refer to ICU stay IDs while others use hospital stay IDs, which slightly complicates data retrieval for a given ID system. Additionally, table layouts vary (*long* versus *wide* data arrangement) and data organization in general is far from consistent over datasets.

1.1. Quick start and organization

We give a quick outline of how **ricu** can be used and how the manuscript is organized.

(i) Download and install the **ricu** package:

- install **ricu** via CRAN with `install.packages("ricu")`
- alternatively, to install from Github use `remotes::install_github("eth-mds/ricu")`

(ii) Get access to datasets and run setup:

- get access to MIMIC-III, eICU and HiRID via Physionet
- get access to AUMCdb via Amsterdam Medical Data Science Website
- set ‘`RICU_PHYSIONET_USER`‘, ‘`RICU_PHYSIONET_PASS`‘ and ‘`RICU_AUMC_TOKEN`‘ environment variables using the obtained information
- run `setup_data_src()` and wait for data download and conversion

A more detailed description of the datasets and the setup process is given in Section 2. There is also a separate vignette to help the user get started.

(iii) Use the main `load_concepts()` function to load 100+ available concepts:

- for example, run `load_concepts(c("age", "glu"), "mimic")` to obtain data on glucose and age concepts in the MIMIC-III dataset
- see `?load_concepts` for more details

A more detailed description of available concepts and data loading is discussed in Section 3.

(iv) Add your own concepts:

- although some commonly appearing concepts are already available (such as demographics, vitals and laboratory values), the user might be interested in adding their own, custom concepts

Addition of custom concepts is possible in two different ways, discussed in detail in Section 3.1.

(v) Add your own dataset:

- the user might be interested in analyzing their own dataset, or using it jointly with the four large in-built datasets in **ricu**
- **ricu** supports the addition of user-specified datasets

Full detail about adding a custom dataset is described in Section 2.3.

Lastly, in Section 4 we briefly show how **ricu** can be used in practice to address clinical questions.

2. Data sources

In order to make data available from different data sources, **ricu** provides abstractions using JSON-formatted configuration files and a set of S3 classes with associated S3 generic functions. This system is designed with extensibility in mind, allowing for incorporation of a wide variety of datasets. Provisions for several large-scale publicly available datasets in terms of required configuration information alongside class-specific implementations of the needed S3 generic functions are part of **ricu**, opening up access to these datasets. Data itself, however, is not part of **ricu** but rather can be downloaded from the Internet using tools provided by **ricu**. While the datasets are publicly available, access has to be granted by the dataset creators individually. Three datasets, MIMIC-III, eICU and HiRID are hosted on PhysioNet (Goldberger, Amaral, Glass, Hausdorff, Ivanov, Mark, Mietus, Moody, Peng, and Stanley 2000), access to which requires an account, while the fourth, AmsterdamUMCdb is currently distributed via a separate platform, requiring a download link.

For both MIMIC-III and eICU, small subsets of data are available as demo datasets that do not require credentialed access to PhysioNet. As the terms for distribution of these demo datasets are less restrictive, they can be made available as data packages **mimic.demo** and **eicu.demo**. Due to size constraints, however they are not available via CRAN, but can be installed from Github as

```
R> install.packages(
+   c("mimic.demo", "eicu.demo"),
+   repos = "https://eth-mds.github.io/physionet-demo"
+ )
```

Provisions for datasets configured to be attached during package loading are made irrespectively of whether data is actually available. Upon first access of a dataset with some

data missing, the user is asked for permission to download in interactive sessions and an error is thrown otherwise. Credentials can either be provided as environment variables (`RICU_PHYSIONET_USER` and `RICU_PHYSIONET_PASS` for access to PhysioNet data, as well as `RICU_AUMC_TOKEN` for AmsterdamUMCdb) and if the corresponding variables are unset, user input is again required in interactive sessions. For non-interactive sessions, functionality is exported such that data can be downloaded and set up ahead of first access (see `?setup_src_data`), which requires the environment variables to be set or credentials to be passed as function arguments.

2.1. Ready to use datasets

Contingent on being granted access by the data owners, several large-scale ICU datasets collected from multiple hospitals in the US and Europe can be set up for access using **ricu** with minimal user effort. Download requires a stable Internet connection as the amount of downloaded data ranges from 5 to 10 GB per dataset, as well as 50 to 100 GB of temporary disk storage for unpacking and preparing the data for efficient access. In terms of permanent storage, again 5 to 10 GB per dataset are required. Despite uncompressed sizes of up to 100 GB for some of the larger tables, memory requirements easily permit importing (and working with) all available tables using Laptop class hardware as only subsets of rows are read at once.

The following paragraphs serve to give quick introductions to the included datasets to offer some guidance, outlining some strengths and weaknesses of each of the datasets. Especially the PhysioNet datasets MIMIC-III and eICU offer good documentation on the respective websites. This section is concluded with a table summarizing similarities and differences among the datasets, outlined in the following paragraphs (see table **??**).

MIMIC-III

The Medical Information Mart for Intensive Care III (MIMIC-III) represents the third iteration of the arguably most influential initiative for collecting and providing to the public large-scale ICU data[2]. The dataset comprises de-identified health related data of roughly 46,000 patients admitted to critical care units of BIDMC during the years 2001-2012. Amounting to just over 61,000 individual ICU admission, data is available on demographics, routine vital sign measurements (at approximately 1 hour resolution), laboratory tests, medication, as well as critical care procedures, organized as a 26-table relational structure.

```
R> mimic
```

One thing of note from a data-organizational perspective is that a change in electronic health care systems occurred in 2008. Owing to this, roughly 48,000 ICU admissions spanning

---

[2]The initial MIMIC (at the time short for Multi-parameter Intelligent Monitoring for Intensive Care) data release dates back 20 years and contained data on roughly 100 patients recorded from patient monitors in the medical, surgical, and cardiac intensive care units of Boston's Beth Israel Hospital during the years 1992-1999 (Moody and Mark 1996). Significantly broadened in scope, MIMIC-II was released 10 years after, now including data on almost 27,000 adult hospital admissions collected from ICUs of Beth Israel Deaconess Medical Center (BIDMC) from 2001 to 2008 (Lee, Scott, Villarroel, Clifford, Saeed, and Mark 2011). Following MIMIC-III, release of MIMIC-IV is imminent with a first development version having been released in summer 2020. This iteration of MIMIC too is planned to be included with **ricu** as soon a first stable version is released.

the years 2001 though 2008 are documented using the CareVue system, while for 2008 and onwards, data was extracted from the MetaVision system. Item identifiers differ between the two systems, requiring queries to consider both ID mappings (heart rate for example being available both as `itemid` number `211` for CareVue and `220045` for MetaVision) as does documentation of infusions and other procedures that are considered as input events (c.f. `inputevents_cv` and `inputevents_mv` tables). Especially with respect to such input event data, MetaVision data generally is of superior quality.

In terms of patient identifiers, MIMIC-III allows for identifying both individual patients (`subject_id`) across hospital admissions (`hadm_id`) and for connecting ICU (re-)admissions (`icustay_id`) to hospital admissions. Using the respective one-to-many relationships, **ricu** can retrieve patient data using any of the above IDs, irrespective of how the raw data is organized.

eICU

Unlike the single-center focus of other datasets, the eICU Collaborative Research Database constitutes an amalgamation of data from critical care units of over 200 hospitals throughout the continental United States. Large-scale data collected via the Philips eICU program which provides telehealth infrastructure for intensive care units, is available from the Philips eICU Research Institute (eRI), albeit neither publicly nor freely. Only data corresponding to roughly 200,000 ICU admissions, sampled from a larger population of over 3 million ICU admissions and stratified by hospital, is being made available via PhysioNet. Patients with discharge dates in 2014 or 2015 were considered, with stays in low acuity units being removed.

```
R> eicu
```

The data is organized into 31 tables and includes patient demographics, routine vital signs, laboratory measurements, medication administrations, admission diagnoses, as well as treatment information. Owing to the wide range of hospitals participating in this data collection initiative, spanning small, rural, non-teaching health centers with fewer than 100 beds to large teaching hospitals with an excess of 500 beds, data availability varies. Even if data was being recorded at the bedside it might end up missing from the eICU dataset due to technical limitations of the collection process. As for patient identifiers, while it is possible to link ICU admissions corresponding to the same hospital stay, it is not possible to identify patients across hospital stays.

Data resolution again varies considerably over included variables. The `vitalperiodic` table stands out as one of the few examples of a *wide* table organization (laying out variables as columns), as opposed to the *long* presentation (following an entity–attribute–value) of most other tables containing patient measurement data. The average time step in `vitalperiodic` is around 5 minutes, but data missingness ranges from around 1% for heart rate and pulse oximetry to around 80-90% for blood pressure measurements, therefore giving approximately hourly resolution for such variables.

HiRID

Developed for early prediction of circulatory failure (Hyland, Faltys, Hüser, Lyu, Gumbsch, Esteban, Bock, Horn, Moor, Rieck, Zimmermann, Bodenham, Borgwardt, Rätsch, and

Merz 2020), the High Time Resolution ICU Dataset (HiRID) contains data on almost 34,000 admissions to the Department of Intensive Care Medicine of the Bern University Hospital, Switzerland, an interdisciplinary 60-bed unit. Given the clear focus on a concrete application during data collection, this dataset is the most limited in terms of breadth of available information, which is also reflected in a comparatively simple data layout comprising only 5 tables[3].

```
R> hirid
```

Collected during the period of January 2008 through June 2016, roughly 700 distinct variables covering routine vital signs, diagnostic test results and treatment parameters are available with variables monitored at the bedside being recorded with two minute time resolution. In terms of demographic information and patient identifier systems however, the data is limited. It is not possible to identify ICU admissions corresponding to individual patients and apart from patient age, sex, weight and height, very little information is available to characterize patients. There is no medical history, no admission diagnoses, only in-ICU mortality information, no unstructured patient data and no information on patient discharge. Furthermore, data on body fluid sampling has been omitted, complicating for example the construction of a Sepsis-3 label (Singer *et al.* 2016).

AmsterdamUMCdb

As a second European dataset, also focusing on increased time-resolution over the US datasets, AmsterdamUMCdb has been made available in late 2019, containing data on over 23,000 intensive care unit and high dependency unit admissions of adult patients during the years 2003 through 2016. The department of Intensive Care at Amsterdam University Medical Center is a mixed medical-surgical ICU with up to 32 bed ICU and 12 bed high dependency units with an average of 1000-2000 yearly admissions. Covering middle ground between the US datasets and HiRID in terms of breadth of included data, while providing a maximal time-resolution of 1 minute, AmsterdamUMCdb constitutes a well organized high quality ICU data resource organized succinctly as a 7-table relational structure.

```
R> aumc
```

A slightly different approach to data anonymization was chosen for this dataset, yielding demographic information such as patient weight, height and age only available as binned variables instead of raw numeric values. Apart from this, there is information on patient origin, mortality, admission diagnoses, as well as numerical measurements including vital parameters, lab results, outputs from drains and catheters, information on administered medication, and other medical procedures. In terms of patient identifiers, it is possible to link ICU admissions corresponding to the same individual, but it is not possible to identify separate hospital admissions.

---

[3]The data is available in three states: as raw data and in preprocessed form, with preprocessed data being represented by two intermediary pipeline stages from (Hyland *et al.* 2020). While **ricu** focuses exclusively on raw data, the *merged* stage represents a selection of variables that were deemed most predictive for determining circulatory failure, which are then merged into 18 meta-variables, representing different clinical concepts. Time stamps in *merged* data are left unchanged, yielding irregular time series, whereas for the *imputed* stage, data is down-sampled to a 5 minute grid and missing values are imputed using a scheme discussed in (Hyland *et al.* 2020).

2.2. Implementation details

Every dataset is represented by an environment with class attributes and associated metadata objects stored as object attributes to that environment. Dataset environments all inherit from `src_env` and from any number of class names constructed from data source name(s) with a suffix `_env` attached. The environment representing MIMIC-III, for example inherits from `src_env` and `mimic_env`, while the corresponding demo dataset inherits from `src_env`, `mimic_env` and `mimic_demo_env`. These sub-classes are later used for adapting the process of data loading to particularities of individual datasets.

A `src_env` contains an active binding per contained table, which returns a `src_tbl` object representing the requested table after having checked that the required data is locally available (and querying the user for download if not). As is the case for `src_env` objects, `src_tbl` objects inherit from additional classes such that certain per-dataset behavior can be customized. The `admissions` table of the MIMIC-III demo dataset for example inherits from `mimic_demo_tbl` and `mimic_tbl` (alongside classes `src_tbl` and `prt`).

```
R> mimic_demo$admissions
```

```
# <mimic_tbl>: [129 x 19]
# ID options:  subject_id (patient) < hadm_id (hadm) < icustay_id (icustay)
# Defaults:    'admission_type' (val)
# Time vars:   'admittime', 'dischtime', 'deathtime', 'edregtime', 'edouttime'
    row_id subject_id hadm_id admittime           dischtime
     <int>      <int>   <int> <dttm>              <dttm>
  1  12258      10006  142345 2164-10-23 21:09:00 2164-11-01 17:15:00
  2  12263      10011  105331 2126-08-14 22:32:00 2126-08-28 18:59:00
  3  12265      10013  165520 2125-10-04 23:36:00 2125-10-07 15:13:00
  4  12269      10017  199207 2149-05-26 17:19:00 2149-06-03 18:42:00
  5  12270      10019  177759 2163-05-14 20:43:00 2163-05-15 12:00:00
...
125  41055      44083  198330 2112-05-28 15:45:00 2112-06-07 16:50:00
126  41070      44154  174245 2178-05-14 20:29:00 2178-05-15 09:45:00
127  41087      44212  163189 2123-11-24 14:14:00 2123-12-30 14:31:00
128  41090      44222  192189 2180-07-19 06:55:00 2180-07-20 13:00:00
129  41092      44228  103379 2170-12-15 03:14:00 2170-12-24 18:00:00
# ... with 119 more rows, and 14 more variables: deathtime <dttm>,
#   admission_type <chr>, admission_location <chr>, discharge_location <chr>,
#   insurance <chr>, language <chr>, religion <chr>, marital_status <chr>,
#   ethnicity <chr>, edregtime <dttm>, edouttime <dttm>, diagnosis <chr>,
#   hospital_expire_flag <int>, has_chartevents_data <int>
```

Powered by the **prt** (Bennett 2021) package, `src_tbl` objects represent tabular data and are stored as a singleton or row-partitioned tables into multiple binary files created by the **fst** (Klik 2020) package. In addition to standard subsetting, `prt` objects can be subsetted via the base R S3 generic function `subset()` and using non-standard evaluation:

```
R> subset(mimic_demo$admissions, subject_id > 44000, language:ethnicity)
```

```
      language           religion marital_status                 ethnicity
1:        ENGL           CATHOLIC         SINGLE                     WHITE
2:        ENGL           CATHOLIC         SINGLE                     WHITE
3:        ENGL           CATHOLIC         SINGLE                     WHITE
4:        ENGL PROTESTANT QUAKER        MARRIED                     WHITE
5:        ENGL       UNOBTAINABLE         SINGLE BLACK/AFRICAN AMERICAN
6:        ENGL           CATHOLIC         SINGLE                     WHITE
7:        ENGL      NOT SPECIFIED         SINGLE                     WHITE
```

This syntax makes it possible to read row-subsets of *long* tables into memory with little memory overhead. While terseness of such an API does introduce potential ambiguity, this is mostly overcome by using the tidy eval framework provided by **rlang** (Henry and Wickham 2020):

```
R> subject_id <- 44000:45000
R> subset(mimic_demo$admissions, .data$subject_id %in% .env$subject_id,
+        subject_id:dischtime)
```

```
   subject_id hadm_id           admittime            dischtime
1:      44083  125157 2112-05-04 08:00:00 2112-05-11 14:15:00
2:      44083  131048 2112-05-22 15:37:00 2112-05-25 13:30:00
3:      44083  198330 2112-05-28 15:45:00 2112-06-07 16:50:00
4:      44154  174245 2178-05-14 20:29:00 2178-05-15 09:45:00
5:      44212  163189 2123-11-24 14:14:00 2123-12-30 14:31:00
6:      44222  192189 2180-07-19 06:55:00 2180-07-20 13:00:00
7:      44228  103379 2170-12-15 03:14:00 2170-12-24 18:00:00
```

By using so-called pronouns (`.data` and `.env`), the distinction can readily be made between a name referring to an object within the context of the data and an object within the context of the calling environment.

Data source set-up

In order to make a dataset accessible to **ricu**, three steps are necessary, each handled by an exported S3 generic function: `download_scr()`, `import_src()` and `attach_src()`. While the first two steps, download and import are one-time procedures, attaching is carried out every time the package namespace is loaded. By default, all data sources known to **ricu** are configured to be attached and in case some data is missing for a given data source, the missing data is downloaded and imported on first access (and contingent on user consent). For downloads to run through without user interaction, several environment variables can be configured:

- `RICU_PHYSIONET_USER`/`RICU_PHYSIONET_PASS`: PhysioNet user name and password with access to the requested dataset.
- `RICU_AUMC_TOKEN`: Download token, extracted from the download URL received when requesting data access.

If any of the required access credentials are not available as environment variables, the user is queried in interactive sessions. Each of the datasets requires 5-10 GB disk space for permanent storage. In addition to that, 50-100 GB of temporary disk storage is required during download and import of any of the datasets, roughly corresponding to the respective maximum table sizes (uncompressed). Memory requirements are kept low by performing all set-up operations only on subsets of rows at the time, such that Laptop-class hardware is sufficient for handling such datasets. Initial data source set up (depending on available download speeds and CPU/disk type) may take upwards of an hour per dataset.

Further environment variables can be set to customize certain aspects of **ricu** data handling:

- `RICU_DATA_PATH`: Data storage location (can be queried by calling `data_dir()`).
- `RICU_CONFIG_PATH`: Comma-separated paths to directories containing configuration files (in addition to the default location; retrievable using `config_paths()`).
- `RICU_SRC_LOAD`: Comma-separated data source names that are set up for being automatically attached on namespace loading (the current set of data sources is available as `auto_attach_srcs()`).

After successful data download, importing prepares tables for efficient random row-access, for which the raw data format (.csv) is not well suited. Tables are read in using **readr** (Wickham and Hester 2020), potentially partitioned row-wise, and re-saved using **fst**. Finally, attaching a dataset creates a corresponding `src_env` object which together with associated meta-data is used by **ricu** to run queries against the data.

Data loading

The lowest level of data access is direct subsetting of `src_tbl` objects as shown at the start of this section (2.2). Building on that, several S3 generic functions successively homogenize data representations, starting with `load_src()`, which provides a string-based interface to `subset()` for all but the row-subsetting expression.

```
R> load_src("admissions", "mimic_demo", subject_id > 44000,
+          cols = c("hadm_id", "admittime", "dischtime"))

    hadm_id           admittime             dischtime
1:   125157 2112-05-04 08:00:00 2112-05-11 14:15:00
2:   131048 2112-05-22 15:37:00 2112-05-25 13:30:00
3:   198330 2112-05-28 15:45:00 2112-06-07 16:50:00
4:   174245 2178-05-14 20:29:00 2178-05-15 09:45:00
5:   163189 2123-11-24 14:14:00 2123-12-30 14:31:00
6:   192189 2180-07-19 06:55:00 2180-07-20 13:00:00
7:   103379 2170-12-15 03:14:00 2170-12-24 18:00:00
```

As data sources differ in their representation of time-stamps, a next step in data homogenization is to converge to a common format: the time difference to the origin time-point of a given ID system.

```
R> load_difftime("admissions", "mimic_demo", subject_id > 44000,
+                cols = c("hadm_id", "admittime", "dischtime"))
```

```
# An 'id_tbl': 7 x 3
# Id var:       'hadm_id'
  hadm_id admittime dischtime
    <int> <drtn>    <drtn>
1  103379 0 mins     13846 mins
2  125157 0 mins     10455 mins
3  131048 0 mins      4193 mins
4  163189 0 mins     51857 mins
5  174245 0 mins       796 mins
6  192189 0 mins      1805 mins
7  198330 0 mins     14465 mins
```

The function `load_difftime()` is expected to return timestamps as base R `difftime` vectors (using `mins` as time unit). The argument `id_hint` can be used to specify a preferred ID system but if not available in raw data, `load_difftime()` will return data using the ID system with highest cardinality. In the above example, if `icustay_id` were requested, data would be returned using `hadm_id`, whereas the a `subject_id` request would be honored, as these two ID columns are available for the `admissions` table.

Building on `load_difftime()` functionality, `load_id()` (and analogously `load_ts()`) returns an `id_tbl` (or `ts_tbl`) object with the requested ID system (passed as `id_var` argument). This uses raw data IDs if available or calls `change_id()` in order to convert to the desired ID system. Similarly, where `load_difftime()` returns data with fixed time interval of one minute, `load_id()` allows for arbitrary time intervals (using `change_interval()`).

```
R> load_id("admissions", "mimic_demo", subject_id > 44000,
+          cols = c("admittime", "dischtime"), id_var = "hadm_id")
```

```
# An 'id_tbl': 7 x 3
# Id var:       'hadm_id'
  hadm_id admittime dischtime
    <int> <drtn>    <drtn>
1  103379 0 hours    230 hours
2  125157 0 hours    174 hours
3  131048 0 hours     69 hours
4  163189 0 hours    864 hours
5  174245 0 hours     13 hours
6  192189 0 hours     30 hours
7  198330 0 hours    241 hours
```

Data source configuration

Data source environments (and corresponding `src_tbl` objects) are constructed using source configuration objects: list-based structures, inheriting from `src_cfg` and from any number of data source-specific class names with suffix `_cfg` appended (as discussed at the beginning of Section 2.2). The exported function `load_src_cfg()` reads a JSON formatted file using **json-lite** (Ooms 2014), and creates a `src_cfg` object per datasource and further therein contained objects.

```
R> cfg <- load_src_cfg("mimic_demo")
R> str(cfg, max.level = 2L, width = 70L)


List of 1
 $ mimic_demo:List of 6
  ..$ name   : chr "mimic_demo"
  ..$ prefix : chr [1:2] "mimic_demo" "mimic"
  ..$ id_cfg : id_cfg [1:3] ‘subject_id‘, ‘hadm_id‘, ‘icustay_id‘
  ..$ col_cfg: col_cfg [1:25] [0, 0, 5, 0, 1], [0, 1, 6, 0, 1], [1, 0, 0...
  ..$ tbl_cfg: tbl_cfg [1:25] [?? x 19; 1], [?? x 24; 1], [?? x 4; 1], [...
  ..$ extra  :List of 1
  ..- attr(*, "class")= chr [1:3] "mimic_demo_cfg" "mimic_cfg" "src_cfg"


R> mi_cfg <- cfg[["mimic_demo"]]
```

In addition to required fields `name` and `prefix` (used as class prefix), as well as further arbitrary fields (`url` in this case), several additional configuration objects are part of `src_cfg`: `id_cfg`, `col_cfg` and `tbl_cfg`.

ID configuration An `id_cfg` object contains an ordered set of key-value pairs representing patient ID systems in a dataset. An implicit assumption currently is that a given patient ID system is used consistently throughout a dataset, meaning that for example an ICU stay ID is always referred to by the same name throughout all tables containing a corresponding column. Owing to the relational origins of these datasets this has been fulfilled in all instances encountered so far. In MIMIC-III, ID systems

```
R> as_id_cfg(mi_cfg)


<id_cfg<mimic_demo[patient < hadm < icustay]>[3]>
     patient          hadm       icustay
‘subject_id‘     ‘hadm_id‘  ‘icustay_id‘
```

are available, allowing for identification of individual patients, their (potentially multiple) hospital admissions over the course of the years and their corresponding ICU admissions (as well as potential re-admissions). Ordering corresponds to cardinality: moving to larger values implies moving along a one-to-many relationship. This information is used in data-loading, whenever the target ID system is not contained in the raw data.

Default column configuration Again used in data loading, this per-table set of key-value pairs specifies column defaults as `col_cfg` object. Each key describes a type of column with special meaning and the corresponding value specifies said column for a given table.

```
R> as_col_cfg(mi_cfg)
```

```
<col_cfg<mimic_demo[id_var, index_var, time_vars, unit_var, val_var]>[25]>
        admissions              callout          caregivers          chartevents
    [0, 0, 5, 0, 1]      [0, 1, 6, 0, 1]      [1, 0, 0, 0, 1]      [0, 1, 2, 1, 1]
          cptevents                d_cpt      d_icd_diagnoses      d_icd_procedures
    [0, 1, 1, 0, 1]      [1, 0, 0, 0, 1]      [1, 0, 0, 0, 1]      [1, 0, 0, 0, 1]
            d_items           d_labitems       datetimeevents         diagnoses_icd
    [1, 0, 0, 0, 1]      [1, 0, 0, 0, 1]      [0, 1, 3, 0, 1]      [0, 0, 0, 0, 1]
           drgcodes             icustays        inputevents_cv        inputevents_mv
    [0, 0, 0, 0, 1]      [0, 1, 2, 0, 1]      [0, 1, 2, 1, 1]      [0, 1, 4, 1, 1]
          labevents microbiologyevents          outputevents              patients
    [0, 1, 1, 1, 1]      [0, 1, 2, 0, 1]      [0, 1, 2, 1, 1]      [0, 0, 4, 0, 1]
       prescriptions procedureevents_mv        procedures_icd              services
    [0, 1, 2, 1, 1]      [0, 1, 4, 1, 1]      [0, 0, 0, 0, 1]      [0, 1, 1, 0, 1]
          transfers
    [0, 1, 2, 0, 1]
```

The following column defaults are currently in use throughout **ricu** but the set of keys can be extended to arbitrary new values:

- `id_var`: In case a table does not contain at least one ID column corresponding to one of the ID systems specified as `id_cfg`, the default ID column can be set on a per-table basis as `id_var`[4].
- `index_var`: A column that is used to define an ordering in time over rows, thereby providing a time-series index[5].
- `time_vars`: Columns which will be treated as time variables (important for converting between ID systems for example), but not as time-series indices[6].
- `unit_var`: Used in concept loading (more specifically for `num_cncpt` concepts, see Section 3.1) to identify columns that represent unit of measurement information.
- `val_var`: Again used when loading data concepts, this identified a default value variable in a table, representing the column of interest to be used as returned data column.

While `id_var`, `index_var` and `time_vars` are used to provide sensible defaults to functions used for general data loading (Section 2.2.2), `unit_var`, `val_var`, as well as potential user-defined defaults are only used in concept loading (see Section 3.3) and therefore need not be prioritized when integrating new data sources until data concepts have been mapped.

Table configuration Finally, `tbl_cfg` objects are used during the initial set-up of a data source. In order to create a representation of a table that is accessible from **ricu** from raw data, several key pieces of information are required:

---

[4]This for example is the case for the `d_items` table in MIMIC-III, which does not contain any patient related data, but holds information on items encoding types of measurements, procedures, etc., used throughout other tables holding actual patient data for identifying the type data point, in line with the relational structure of the data source.

[5]For the MIMIC-III table `inputevents_mv`, one of the four available time variables lends itself to be used as index variable more than the other candidates and therefore is set as default.

[6]In case of the `admissions` table in MIMIC-III for example, a total of five columns are considered to be time variables, none of which stands out as potential `index_var`.

- File name(s): In the simplest case, a single file corresponds to a single table. Other scenarios that have been encountered (and are therefore handled) include tables partitioned into multiple files and .tar archives containing multiple tables.

- Column specification: For each column, the expected data type has to be known, as well as a pair of names, one corresponding to the raw data column name and one corresponding to the column name to be used within **ricu**.

- (Optional) number of rows: Used as sanity check whenever available.

- (Optional) partitioning information: For very *long* tables it can be useful to specify a row-partitioning. This currently is only possible by applying a vector of breakpoints to a single numeric column, thereby defining a grouping.

```
R> as_tbl_cfg(mi_cfg)
```

```
<tbl_cfg<mimic_demo[rows x cols; partitions]>[25]>
       admissions            callout          caregivers          chartevents
      [?? x 19; 1]       [?? x 24; 1]        [?? x 4; 1]         [?? x 15; 2]
        cptevents              d_cpt     d_icd_diagnoses     d_icd_procedures
      [?? x 12; 1]        [?? x 9; 1]        [?? x 4; 1]          [?? x 4; 1]
          d_items         d_labitems      datetimeevents        diagnoses_icd
      [?? x 10; 1]        [?? x 6; 1]       [?? x 14; 1]          [?? x 5; 1]
         drgcodes           icustays       inputevents_cv       inputevents_mv
       [?? x 8; 1]       [?? x 12; 1]       [?? x 22; 1]         [?? x 31; 1]
        labevents microbiologyevents         outputevents             patients
       [?? x 9; 1]       [?? x 16; 1]       [?? x 13; 1]          [?? x 8; 1]
    prescriptions procedureevents_mv       procedures_icd             services
      [?? x 19; 1]       [?? x 25; 1]        [?? x 5; 1]          [?? x 6; 1]
        transfers
      [?? x 13; 1]
```

For the `chartevents` table of the MIMIC-III demo dataset, for example, rows are partitioned into two groups, while all other tables are represented by a single partition. Furthermore, the expected number of rows is unknown (`??`) as this is missing from the corresponding `tbl_cfg` object.

2.3. Adding external datasets

In order to add a new dataset to **ricu**, several aspects outlined in the previous subsections require consideration. For illustration purposes, code for integrating AmsterdamUMCdb as external dataset is available from Github. While this is no longer needed for using the `aumc` data source, the repository will remain as it might serve as template to integration of new datasets.

Using a configuration file as described in Section 2.2.3 and pointing the `RICU_CONFIG_PATH` environment variable at its location, data can be prepared for use with **ricu** using `import_src()` and made available to **ricu** using `attach_src()`. In addition to providing a configuration file, dataset-specific implementations of some of the S3 generic functions involved in data-loading

might be required. In the case of AmsterdamUMCdb, a class-specific implementation of `load_difftime()` is required, as raw time-stamps are recoded in milliseconds (instead of minutes), as well as a class-specific implementation of the S3 generic function `id_win_helper()`[7].

3. Data concepts

One of the key components of **ricu** is a scheme for specifying how to retrieve data corresponding to pre-defined clinical concepts from a given data source, in turn enabling dataset agnostic code for analysis. Heart rate, for example can be loaded using the `hr` concept as

```
R> load_concepts("hr", c("mimic_demo", "eicu_demo"), verbose = FALSE)
```

```
# A 'ts_tbl': 152,197 x 4
# Id vars:    'source', 'icustay_id'
# Units:      'hr' [bpm]
# Index var:  'charttime' (1 hours)
        source    icustay_id charttime    hr
        <chr>          <int> <drtn>     <dbl>
      1 eicu_demo     141764   0 hours    119
      2 eicu_demo     141764   1 hours    105
      3 eicu_demo     141764   2 hours     95
      4 eicu_demo     141764   3 hours    106
      5 eicu_demo     141764   4 hours    102
    ...
152,193 mimic_demo    298685 314 hours     60
152,194 mimic_demo    298685 315 hours     56
152,195 mimic_demo    298685 316 hours     50
152,196 mimic_demo    298685 317 hours     48
152,197 mimic_demo    298685 318 hours      0
# ... with 152,187 more rows
```

This requires some form of infrastructure for concisely specifying how to retrieve data subsets (Section 3.1), which is both extensible (to new concepts and new datasets) and flexible enough to handle concept-specific pre-processing. Additionally, **ricu** has included a dictionary with over 100 concepts implemented for all four supported datasets (where possible; see also Section 3.3). A quick remark on terminology before diving into more details on how to specify data concepts: A *concept* corresponds to a clinical variable such as a bilirubin measurement or the ventilation status of a patient, and an *item* encodes how to retrieve data corresponding to a given concept from a data source. A *concept* therefore contains several *items* (zero, one or several are possible per data source).

---

[7]This method is used as part of `change_id()` which is called whenever the requested ID system is not available in raw data. A component central to conversion between patient ID systems is a table which contains patient IDs as columns alongside columns with start and end-points for each, thereby specifying a mapping between ID systems. Dataset-specific construction of such a table is handled by the S3 generic function `id_win_helper()`. As construction of such tables can be expensive (involving several merge operations of tables with 10^4-10^5 rows) and used frequently (potentially with every single data request), the resulting table is cached in memory with session persistence.

3.1. Concept specification

Similarly to data source configuration (discussed in Section 2.2.3), concept specification relies on JSON-formatted text files. A default dictionary of concepts is included with **ricu** containing a selection of commonly used clinical concepts. Several types of concepts exist within **ricu** and with extensibility in mind, new types can easily be added.

All concepts consist of minimal meta-data including a name, target class (defaults to `ts_tbl`; see Section 3.2), an aggregation specification[8] and class information (defaults to `num_concept`), as well as optional `description` and `category` information. Adding to that, depending on concept class, further fields can be added. In the case of the most widespread concept type (`num_cncpt`; used to represent numeric data) this is `unit` which encodes one (or several synonymous) unit(s) of measurement, as well as a minimal and maximal plausible values (specified as `min` and `max`). The concept for heart rate data (`hr`) for example can be specified as

```
{
  "hr": {
    "unit": ["bpm", "/min"],
    "min": 0,
    "max": 300,
    "description": "heart rate",
    "category": "routine vital signs",
    "sources": {
       ...
    }
  }
}
```

Meta-data is used during concept loading for data-preprocessing. For numeric concepts, the specified measurement unit is compared to that of the data (if available), with messages being displayed in case of mismatches, while the range of plausible values is used to filter out measurements that fall outside the specified interval. Other types of concepts include categorical concepts (`fct_cncpt`), concept representing binary data (`lgl_cncpt`), as well as recursive concepts (`rec_cncpt`), which build on other *atomic* concepts[9].

---

[8]Every concept needs a default aggregation method which can be used during data loading to return data that is unique per key (either per `id_vars` group or per combination of `id_vars` and `index_var`) otherwise down-stream merging of multiple concepts is ill-defined. The aggregation default can be overridden during loading or as specification of a `rec_cncpt` object. If no aggregation method is explicitly indicated the global default is `first()` for character, `median()` for numeric and `sum()` for logical vectors. For logical data, if a concept of type `lgl_cncpt` is used, the count of `TRUE` values is converted back to logical, thereby providing `any()` type functionality.

[9]An example for a recursive concept is the PaO_2/FiO_2 ratio, used for instance to assess patients with acute respiratory distress syndrome (ARDS) or for sepsis-related organ failure assessment (SOFA) (Villar, Pérez-Méndez, Blanco, Añón, Blanch, Belda, Santos-Bouza, Fernández, Kacmarek, and Spanish Initiative for Epidemiology and Therapies for ARDS (SIESTA) Network 2013; Vincent *et al.* 1996). Given both PaO_2 and FiO_2 as individual concepts, the PaO_2/FiO_2 ratio is provided by **ricu** as a recursive concept (`pafi`), requesting the two atomic concepts `pao2` and `fio2` and performing some form of imputation for when at a given time step one or both values are missing.

Specification of how data can be retrieved from a data source is encoded by data *items*. Lists of data items (associated with data source names) are provided as `sources` element (instead of `...` in the above code block). For the demo datasets corresponding eICU and MIMIC-III, heart rate data retrieval is specified as

```
{
  "eicu_demo": [
    {
      "table": "vitalperiodic",
      "val_var": "heartrate",
      "class": "col_itm"
    }
  ],
  "mimic_demo": [
    {
      "ids": [211, 220045],
      "table": "chartevents",
      "sub_var": "itemid"
    }
  ]
}
```

Analogously to how different types of concepts are used to represent different types of data, different types of items handle different types of data loading. The most common scenario is selecting a subset of rows from a table by matching a set of ID values (`sub_itm`). In the above example, heart rate data in MIMIC-III can be located by searching for ID values 211 and 220045 in column `itemid` of table `chartevents` (heart rate data is stored in *long* format). Conversely, heart rate data in eICU is stored in *wide* format, requiring no row-subsetting. Column `heartrate` of table `vitalperiodic` contains all corresponding data and such data situations are handled by the `col_itm` class. Other item classes include `rgx_itm` where a regular expression is used for selecting rows and `fun_itm` where an arbitrary function can be used for data loading. If a data loading scenario is not covered by these classes, adding further `itm` subclasses is encouraged.

In order to extend the current concept library both to new datasets and new concepts, further JSON files can be incorporated by adding their paths to `RICU_CONFIG_PATH`. Concepts with names that already exist are only used for their `sources` entries, such that `hr` for `new_dataset` can be specified as

```
"hr": {
  "sources": {
    "new_dataset": [
      {
        "ids": 6640,
        "table": "numericitems",
        "sub_var": "itemid"
      }
```

```
    ]
  }
}
```

whereas concepts with non-existing names are treated as new concepts.

Central to providing the required flexibility for loading of certain data concepts that require some specific pre-processing are callback functions that can be specified for several *item* types. Functions (with appropriate signatures), designated as `callback` functions, are invoked on individual data items, before concept-related preprocessing is applied. A common scenario for this is unit of measurement conversion: In MIMIC-III data for example, several `itemid` values correspond to temperature measurements, some of which refer to temperatures measured in degrees Celsius whereas others are used for measurements in degrees Fahrenheit. As the information encoding which measurement corresponds to which `itemid` values is no longer available during concept-related preprocessing, this is best resolved at the level of individual data items. Several function factories are available for generating callback functions and `convert_unit()` is intended for covering unit conversions. Data *items* corresponding to the `temp` concept for MIMIC-III are specified as

```
{
  "mimic_demo": [
    {
      "ids": [676, 677, 223762],
      "table": "chartevents",
      "sub_var": "itemid"
    },
    {
      "ids": [678, 679, 223761, 224027],
      "table": "chartevents",
      "sub_var": "itemid",
      "callback": "convert_unit(fahr_to_cels, 'C', 'f')"
    }
  ]
}
```

indicating that for ID values 676, 677 and 223762 no pre-processing is required and for the remaining ID values the function `fahr_to_cels()` is applied to entries of the `val_var` column where the regular expression `"f"` is `TRUE` for the `unit_var` column (the values of which being ultimately replaced with `"C"`).

3.2. Data classes

In order to represent tabular ICU data, **ricu** provides several classes, all inheriting from `data.table`. The most basic of which, `id_tbl`, marks one (or several) columns as `id_vars` which serve to define a grouping (i.e. identify patients or unit stays). Inheriting from `id_tbl`, `ts_tbl` is capable of representing grouped time-series data. In addition to `id_var` column(s), a single column is marked as `index_var` and is required to hold a base R `difftime` vec-

tor. Furthermore, `ts_tbl` contains a scalar-valued `difftime` object as `interval` attribute, specifying the time-series step size[10].

Meta data is transiently added to `data.table` objects by classes inheriting from `id_tbl` and S3 generic functions which allow for object modifications, down-casting is implicit:

```
R> (dat <- ts_tbl(a = 1:5, b = hours(1:5), c = rnorm(5)))

# A 'ts_tbl': 5 x 3
# Id var:     'a'
# Index var:  'b' (1 hours)
      a b           c
  <int> <drtn>   <dbl>
1     1 1 hours  1.08
2     2 2 hours -0.179
3     3 3 hours -0.300
4     4 4 hours  0.953
5     5 5 hours  1.49

R> dat[["b"]] <- dat[["b"]] + mins(30)
R> dat

# An 'id_tbl': 5 x 3
# Id var:     'a'
      a b             c
  <int> <drtn>     <dbl>
1     1  5400 secs  1.08
2     2  9000 secs -0.179
3     3 12600 secs -0.300
4     4 16200 secs  0.953
5     5 19800 secs  1.49
```

Due to time-series step size of `dat` being specified as 1 hour, an internal inconsistency is encountered when shifting time stamps by 30 minutes, as time-steps are no longer multiples of the time-series interval, in turn causing down-casting to `id_tbl`. If column `a` were to be removed, direct down-casting to `data.table` would be required in order to resolve inconsistencies[11].

Utilizing the attached meta-data, several utility functions can be called with concise semantics. This includes functions for sorting, checking for duplicates, aggregating data per combination of `id_vars` (and time-step), checking time series data for gaps, verifying whether the time-series is regular and converting between irregular and regular time-series, as well

---

[10]As further extension, a `win_tbl` object is being considered for inclusion, capable of representing time intervals. Such an object could prove convenient for example when dealing with infusions as infusion parameters such as medication rate frequently are specified with explicit begin and end times (see Section 3.3.3).

[11]Updating an object inheriting from `id_tbl` using `data.table::set()` bypasses consistency checks as this is not an S3 generic function and therefore its behavior cannot be tailored to requirements of `id_tbl` objects. It therefore is up to the user to avoid vitiating `id_tbl` objects in such a way.

as functions for several types of moving window operations. Adding to those class-specific implementations, `id_tbl` objects inherit from `data.table` (and therefore from `data.frame`), ensuring compatibility with a wide range of functionality targeted at these base-classes.

3.3. Clinical concepts

The current selection of clinical concepts that is included with **ricu** covers many physiological variables that are available throughout the included datasets. Treatment-related information on the other hand, being more heterogeneous in nature and therefore harder to harmonize across datasets, has been added on an as-needed basis and therefore is more limited in breadth.

Available concepts can be enumerated using `load_dictionary()` and the utility function `explain_dictionary()` can be used to display some concept meta-data.

```
R> dict <- load_dictionary(c("mimic_demo", "eicu_demo"))
R> head(dict)


<concept[6]>
                            abx                              adh_rate
          antibiotics <lgl_cncpt[4]>      vasopressin rate <num_cncpt[3]>
                            adm                                   age
patient admission type <fct_cncpt[2]>         patient age <num_cncpt[2]>
                            alb                                   alp
              albumin <num_cncpt[2]>  alkaline phosphatase <num_cncpt[2]>
```

```
R> explain_dictionary(head(dict))


      name      category             description
1      abx   medications             antibiotics
2 adh_rate   medications        vasopressin rate
3      adm  demographics  patient admission type
4      age  demographics             patient age
5      alb     chemistry                 albumin
6      alp     chemistry    alkaline phosphatase
```

The following sub-sections serve to introduce some of the included concepts as well as highlight limitations that come with current implementations. Grouping the available concepts by category yields the following counts

```
R> table(vapply(dict, '[[', character(1L), "category"))


   blood gas    chemistry demographics   hematology  medications microbiology
          10           21            6           20           16            1
neurological      outcome       output  respiratory       vitals
           8           19            2           10            6
```

Physiological data

The largest and most well established group of concepts (covering more than half of all currently included concepts) includes physiological patient measurements such as routine vital signs, respiratory variables, fluid discharge amounts, as well as many kinds of laboratory tests including blood gas measurements, chemical analysis of body fluids and hematology assays.

```
R> load_concepts(c("alb", "glu"), "mimic_demo", interval = mins(15L),
+                verbose = FALSE)

# A 'ts_tbl': 1,965 x 4
# Id var:      'icustay_id'
# Units:       'alb' [g/dL], 'glu' [mg/dL]
# Index var:  'charttime' (15 mins)
      icustay_id charttime    alb   glu
           <int> <drtn>      <dbl> <dbl>
    1     201006 -3495 mins  NA     116
    2     201006 -2745 mins  NA      83
    3     201006 -1275 mins  NA      91
    4     201006    15 mins  2.4    175
    5     201006   675 mins  NA     129
  ...
1,961     298685 15600 mins  NA     159
1,962     298685 16365 mins  2.2    153
1,963     298685 17400 mins  NA     182
1,964     298685 17595 mins  NA     122
1,965     298685 17955 mins  2.5    121
# ... with 1,955 more rows
```

Most concepts of this kind are represented by `num_cncpt` objects with an associated unit of measurement and a range of permissible values. Data is mainly returned as `ts_tbl` objects, representing time-dependent observations. Apart from conversion to a common unit (possibly using the `convert_unit()` callback function), little has to be done in terms of pre-processing: values are simply reported at time-points rounded to the requested interval.

Patient demographics

Moving on from dynamic, time-varying patient data, this group of concepts focuses on static patient information. While the assumption of remaining constant throughout a stay is likely to hold for variables such as patient sex or height this is only approximately true for others including age or weight. Nevertheless such effects are ignored and concepts of this group will be mainly returned as `id_tbl` objects with no corresponding time-stamps included.

Whenever requesting concepts which are returned with associated time-stamps (e.g. glucose) alongside time-constant data (e.g. age), merging will duplicate static data over all time-points.

```
R> load_concepts(c("age", "glu"), "mimic_demo", verbose = FALSE)
```

```
# A 'ts_tbl': 1,914 x 4
# Id var:       'icustay_id'
# Units:        'age' [years], 'glu' [mg/dL]
# Index var: 'charttime' (1 hours)
      icustay_id charttime    age    glu
            <int> <drtn>      <dbl> <dbl>
    1     201006 -58 hours   68.9    116
    2     201006 -45 hours   68.9     83
    3     201006 -21 hours   68.9     91
    4     201006   0 hours   68.9    175
    5     201006  11 hours   68.9    129
  ...
1,910     298685 260 hours   80.1    159
1,911     298685 272 hours   80.1    153
1,912     298685 290 hours   80.1    182
1,913     298685 293 hours   80.1    122
1,914     298685 299 hours   80.1    121
# ... with 1,904 more rows
```

Despite a best-effort approach, data availability can be a limiting factor. While for physiological variables, there is good agreement even across continents, data-privacy considerations, as well as lack of a common standard for data encoding, may cause issues that are hard to resolve. In some cases, this can be somewhat mitigated while in others, this is a limitation to be kept in mind. In AmsterdamUMCdb, for example, patient age, height and weight are not available as continuous variables, but as factor with patients binned into groups. Such variables are then approximated by returning the respective mid-points of groups for `aumc` data[12]. Other concepts, such as `adm` (categorizing admission types) or a prospective `icd` concept (diagnoses as ICD-9 codes) can only return data if available from the data source in question. Unfortunately, neither `aumc` nor `hirid` contain ICD-9 encoded diagnoses, and in the case of `hirid`, no diagnosis information is available at all.

Treatment-related information

The largest group of concepts dealing with treatment-related information is described by the `medications` category. In addition to drug administrations, only basic ventilation information is currently provided as ready to use concept. Just like availability of common ICU procedures, patient medication is also underdeveloped, covering mainly vasopressor administrations, as well as corticosteroids and antibiotics. The current concepts retrieving treatment-related information are mostly focused on providing data required for constructing clinical scores described in Section 3.3.4.

Ventilation is represented by several concepts: a ventilation indicator variable (`vent_ind`), as well as ventilation durations (`vent_dur`) are constructed from start and events (`vent_start` and `vent_end`). This includes any kind of mechanical ventilation (invasive via an endotracheal

---

[12]Prioritizing consistency over accuracy, one could apply the same binning to datasets which report numeric values, but the concepts included with **ricu** attempt to strike a balance between consistency and amount of applied pre-processing. With the extensible architecture of data concepts, however, such categorical variants of patient demographic concepts could easily be added.

or tracheostomy tube), as well as non-invasive ventilation via face or nasal masks. In line with other concepts belonging to this group, the current state is far from being comprehensive and expansion to further ventilation parameters is desirable.

The singular concept addressing antibiotics (`abx`) returns an indicator signaling whenever an antibiotic was administered. This includes any route of administration (intravenous, oral, topical, etc.) and does neither report dosage, nor active ingredient. Finally, vasopressor administration is reported by several concepts representing different vasoactive drugs (including dopamine, dobutamine, epinephrine, noreponephrine and vasopressin), as well as different administration aspects such as rate, duration (and for use in SOFA scoring, rate administered for at least 60 minutes).

```
R> load_concepts(c("abx", "vent_ind", "norepi_rate", "norepi_dur"),
+                "mimic_demo", verbose = FALSE)

# A 'ts_tbl': 12,547 x 6
# Id var:     'icustay_id'
# Units:      'norepi_rate' [mcg/kg/min]
# Index var:  'startdate' (1 hours)
       icustay_id startdate abx   vent_ind norepi_rate norepi_dur
            <int> <drtn>    <lgl> <lgl>          <dbl> <drtn>
     1     201006 -40 hours TRUE  NA                NA NA hours
     2     201006 -16 hours TRUE  NA                NA NA hours
     3     201006   7 hours TRUE  NA                NA NA hours
     4     201006   8 hours NA    TRUE              NA 60 hours
     5     201006   9 hours NA    TRUE          0.0460 NA hours
   ...
12,543     298685 612 hours NA    TRUE              NA NA hours
12,544     298685 613 hours NA    TRUE              NA NA hours
12,545     298685 614 hours NA    TRUE              NA NA hours
12,546     298685 615 hours NA    TRUE              NA NA hours
12,547     298685 616 hours NA    TRUE              NA NA hours
# ... with 12,537 more rows
```

As cautioned in Section 3.3.2, variability in data reporting across datasets can lead to issues: the `prescriptions` table included with MIMIC-III, for example, reports time-stamps as dates only, yielding a discrepancy of up to 24 hours when merged with data where time-accuracy is on the order of minutes. This effect is somewhat mitigated by shifting time-stamps from midnight to mid-day, but the underlying accuracy issue of course remains. Another problem exists with concepts that attempt to report administration windows, as some datasets do not describe infusions with clear cut start/endpoints but rather report infusion parameters at (somewhat) regular time intervals. This can cause artifacts when the requested time step-size deviates from the dataset inherent time grid.

Outcomes

A group of more loosely associated concepts can be used to describe patient state. This includes common clinical endpoints, such as death or length of ICU stay, as well as scoring

systems such as SOFA, the systemic inflammatory response syndrome (SIRS; Bone, Sibbald, and Sprung 1992) criterion, the National Early Warning Score (NEWS; Jones 2012) and the Modified Early Warning Score (MEWS; Subbe, Kruger, Rutherford, and Gemmel 2001).

While the more straightforward outcomes can be retrieved directly from data, clinical scores often incorporate multiple variables, based upon which a numeric score is constructed. This can typically be achieved by using concepts of type `rec_cncpt`, specifying the needed components and supplying a callback function that applies rules for score construction.

```
R> load_concepts(c("sirs", "death"), "mimic_demo", verbose = FALSE,
+              keep_components = TRUE)

# A 'ts_tbl': 14,295 x 8
# Id var:     'icustay_id'
# Index var:  'charttime' (1 hours)
```

	icustay_id	charttime	sirs	death	temp_comp	hr_comp	resp_comp	wbc_comp
	\<int>	\<drtn>	\<dbl>	\<lgl>	\<int>	\<int>	\<int>	\<int>
1	201006	-58 hours	1	NA	NA	NA	NA	1
2	201006	-45 hours	1	NA	NA	NA	NA	1
3	201006	-21 hours	1	NA	NA	NA	NA	1
4	201006	-10 hours	2	NA	NA	NA	1	1
5	201006	0 hours	3	NA	0	1	1	1
...								
14,291	298685	314 hours	2	NA	0	0	1	1
14,292	298685	315 hours	2	NA	0	0	1	1
14,293	298685	316 hours	2	NA	0	0	1	1
14,294	298685	317 hours	1	NA	0	0	0	1
14,295	298685	318 hours	1	TRUE	0	0	NA	1

```
# ... with 14,285 more rows
```

Callback functions can become rather involved (especially for more complex concepts such as SOFA) and may include arbitrary arguments to tune their behavior. As callback functions to `rec_cncpt` objects are typically called internally from `load_concepts()`, arguments not used by `load_concepts()`, such as `keep_components` in the above example (causing not only the score column, but also individual score components to be retained) are forwarded[13].

4. Examples

In order to briefly illustrate how **ricu** could be applied to real-world clinical questions, two toy examples are provided in the following sections. While the first example fully relies on data concepts that are included with **ricu**, the second one explores both how some data pre-processing can be added to an existing concept by creating a new `rec_cncpt` and how to create an new data concept altogether.

---

[13]Some care has to be taken as when requesting multiple concepts within the same call to `load_concepts()` all involved callback functions will be called with the same forwarded arguments. When for example requesting multiple scores, it is currently not possible to enable `keep_components` for only a subset thereof. Furthermore this set-up implies that

4.1. Lactate and mortality

First, we investigate the association of lactate levels and mortality. This problem has been studied before and it is widely accepted that both static and dynamic lactate indices are associated with increased mortality (Haas, Lange, Saugel, Petzoldt, Fuhrmann, Metschke, and Kluge 2016; Nichol, Bailey, Egi, Pettila, French, Stachowski, Reade, Cooper, and Bellomo 2011; Van Beest, Brander, Jansen, Rommes, Kuiper, and Spronk 2013). In order to model this relationship, we fit a time-varying proportional hazards Cox model (Therneau and Grambsch 2000; Therneau and Lumley 2015), which includes the SOFA score as a general predictor of illness severity, using MIMIC-III data. Furthermore, for the sake of this example, we are only interested in patients admitted from 2008 onwards of ages 25 to 65 years old.

```
R> cohort <- load_id("icustays", "mimic", dbsource == "metavision",
+                    cols = NULL)
R> cohort <- load_concepts("age", "mimic", patient_ids = cohort,
+                          verbose = FALSE)
R>
R> dat <- load_concepts(c("lact", "death", "sofa", "sex"), "mimic",
+                       patient_ids = cohort[age > 25 & age < 65],
+                       verbose = FALSE)
R>
R> dat <- dat[, head(.SD, n = match(TRUE, death, .N)), by = c(id_vars(dat))]
R> dat <- fill_gaps(dat)
R>
R> dat <- replace_na(dat, c(NA, FALSE), type = c("locf", "const"),
+                    by_ref = TRUE, vars = c("lact", "death"),
+                    by = id_vars(dat))
R>
R> cox_mod <- coxph(
+   Surv(charttime - 1L, charttime, death) ~ lact + sofa,
+   data = dat
+ )
```

After loading the data, some minor pre-processing is still required before modeling: first, we want to make sure we only use data up to (and including) the hour in which the `death` flag switches to `TRUE`. After that we impute missing values for `lact` using a last observation carry forward (locf) scheme (observing the patient grouping) and we simply replace missing `death` values with the value `FALSE`. The resulting model fit can be visualized as:

A simple exploration already shows that the increased values of lactate are associated with mortality, even after adjusting for the SOFA score.

4.2. Diabetes and insulin treatment

For the next example, again using MIMIC-III data, we turn to the usage of co-morbidities and treatment related information. We look at the amount of insulin administered to patients in the first 24 hours from their ICU admission. In particular, we investigate if diabetic patients receive more insulin in the first day of their stay compared to non-diabetic patients. For

this we create two concepts: `ins24`, a binned variable representing the cumulative amount of insulin administered within the first 24 hours of an ICU admission, and `diab`, a logical variable encoding diabetes co-morbidity.

As there already is an insulin concept available, `ins24` can be implemented as `rec_cncpt`, loading `ins` with aggregation set to `sum()` (instead of `median()`) and inserting the callback function `ins_cb()` into the loading process. The callback function takes care of the preprocessing steps outlined above: first data is subsetted to fall into the the first 24 hours of ICU admissions, followed by binning of summed values.

```
R> ins_breaks <- c(0, 1, 10, 20, 40, Inf)
R>
R> ins_cb <- function(ins, ...) {
+
+   day_one <- function(x) x >= hours(0L) & x <= hours(24L)
+
+   idx_var <- index_var(ins)
+   ids_var <- id_vars(ins)
+
+   ins <- ins[
+     day_one(get(idx_var)), list(ins24 = sum(ins)), by = c(ids_var)
+   ]
+
+   ins <- ins[,
+     ins24 := list(cut(ins24, breaks = ins_breaks, right = FALSE))
+   ]
+
+   ins
+ }
R>
R> ins24 <- load_dictionary("mimic", "ins")
R> ins24 <- concept("ins24", ins24, "insulin in first 24h", aggregate = "sum",
+                   callback = ins_cb, target = "id_tbl", class = "rec_cncpt")
```

The binary diabetes concept can be implemented as `lgl_cncpt`, for which ICD-9 codes are matched using a regular expression. As we're not only interested in retrieving diabetic patients, a `col_itm` is more suited for data retrieval over an `rgx_itm` and for creating the required callback function that produces a logical vector we can use `transform_fun()` coupled with a function like `grep_diab()`. The two concepts are then combined using `c()` and loaded via `load_concepts()`.

```
R> grep_diab <- function(x) grepl("^250\\.?[0-9]{2}$", x)
R>
R> diab  <- item("mimic", table = "diagnoses_icd",
+               callback = transform_fun(grep_diab), class = "col_itm")
R> diab  <- concept("diab", diab, "diabetes", target = "id_tbl",
+                   class = "lgl_cncpt")
```

```
R>
R> dat <- load_concepts(c(ins24, diab), id_type = "icustay", verbose = FALSE)
R> dat <- replace_na(dat, "[0,1)", vars = "ins24")
R> dat
```

After this, we can visualize the difference between the two groups with a histogram:

The plot suggests that during the first day of ICU stay, perhaps unsurprisingly, diabetic patients receive more insulin compared to non-diabetic patients, especially as insulin dose increases.

5. Acknowledgments

References

Adibuzzaman M, Musselman K, Johnson A, Brown P, Pitluk Z, Grama A (2016). "Closing the data loop: An integrated open access analysis platform for the mimic database." In *2016 Computing in Cardiology Conference (CinC)*, pp. 137–140. IEEE.

Bennett N (2021). *prt: Tabular Data Backed by Partitioned 'fst' Files.* R package version 0.1.3, URL https://CRAN.R-project.org/package=prt.

Bone RC, Sibbald WJ, Sprung CL (1992). "The ACCP-SCCM Consensus Conference on Sepsis and Organ Failure." *Chest*, **101**(6), 1481 – 1483. ISSN 0012-3692.

Desautels T, Calvert J, Hoffman J, Jay M, Kerem Y, Shieh L, Shimabukuro D, Chettipally U, Feldman MD, Barton C, *et al.* (2016). "Prediction of sepsis in the intensive care unit with minimal electronic health record data: a machine learning approach." *JMIR medical informatics*, **4**(3), e28.

Evans R (2016). "Electronic health records: then, now, and in the future." *Yearbook of medical informatics*, **25**(S 01), S48–S61.

Faltys M, Zimmermann M, Lyu X, Hüser M, Hyland SL, Rätsch G, Merz TM (2021). "HiRID, a high time-resolution ICU dataset (version 1.1.1)." PhysioNet.

Fleuren LM, Klausch TLT, Zwager CL, Schoonmade LJ, Guo T, Roggeveen LF, Swart EL, Girbes ARJ, Thoral P, Ercole A, Hoogendoorn M, Elbers PWG (2020). "Machine learning for the prediction of sepsis: a systematic review and meta-analysis of diagnostic test accuracy." *Intensive Care Medicine*, **46**(3), 383–400.

Futoma J, Hariharan S, Sendak M, Brajer N, Clement M, Bedoya A, O'Brien C, Heller K (2017). "An improved multi-output gaussian process rnn with real-time validation for early sepsis detection." *arXiv preprint arXiv:1708.05894*.

Goldberger AL, Amaral LAN, Glass L, Hausdorff JM, Ivanov PC, Mark RG, Mietus JE, Moody GB, Peng CK, Stanley HE (2000). "PhysioBank, PhysioToolkit, and PhysioNet." *Circulation*, **101**(23), e215–e220.

Haas SA, Lange T, Saugel B, Petzoldt M, Fuhrmann V, Metschke M, Kluge S (2016). "Severe hyperlactatemia, lactate clearance and mortality in unselected critically ill patients." *Intensive care medicine*, **42**(2), 202–210.

Henry L, Wickham H (2020). *rlang: Functions for Base Types and Core R and 'Tidyverse' Features*. R package version 0.4.9, URL `https://CRAN.R-project.org/package=rlang`.

Hyland SL, Faltys M, Hüser M, Lyu X, Gumbsch T, Esteban C, Bock C, Horn M, Moor M, Rieck B, Zimmermann M, Bodenham D, Borgwardt K, Rätsch G, Merz TM (2020). "Early prediction of circulatory failure in the intensive care unit using machine learning." *Nature Medicine*, **26**(3), 364–373.

Jiang F, Jiang Y, Zhi H, Dong Y, Li H, Ma S, Wang Y, Dong Q, Shen H, Wang Y (2017). "Artificial intelligence in healthcare: past, present and future." *Stroke and vascular neurology*, **2**(4), 230–243.

Johnson AE, Pollard TJ, Shen L, Li-wei HL, Feng M, Ghassemi M, Moody B, Szolovits P, Celi LA, Mark RG (2016). "MIMIC-III, a freely accessible critical care database." *Scientific data*, **3**, 160035.

Johnson AEW, Aboab J, Raffa JD, Pollard TJ, Deliberato RO, Celi LA, Stone DJ (2018). "A Comparative Analysis of Sepsis Identification Methods in an Electronic Database." *Critical care medicine*, **46**(4), 494–499.

Jones M (2012). "NEWSDIG: The National Early Warning Score Development and Implementation Group." *Clinical Medicine*, **12**(6), 501–503. `doi:10.7861/clinmedicine.12-6-501`.

Kam HJ, Kim HY (2017). "Learning representations for the early detection of sepsis with deep neural networks." *Computers in biology and medicine*, **89**, 248–255.

Klik M (2020). *fst: Lightning Fast Serialization of Data Frames*. R package version 0.9.4, URL `https://CRAN.R-project.org/package=fst`.

Lee J, Scott D, Villarroel M, Clifford G, Saeed M, Mark R (2011). "Open-access MIMIC-II database for intensive care research." In *Annual International Conference of the IEEE Engineering in Medicine and Biology Society. IEEE Engineering in Medicine and Biology Society. Conference*, volume 2011, pp. 8315–8.

Moody GB, Mark RG (1996). "A database to support development and evaluation of intelligent intensive care monitoring." In *Computers in Cardiology*, pp. 657–660.

Nemati S, Holder A, Razmi F, Stanley MD, Clifford GD, Buchman TG (2018). "An interpretable machine learning model for accurate prediction of sepsis in the ICU." *Critical care medicine*, **46**(4), 547–553.

Nichol A, Bailey M, Egi M, Pettila V, French C, Stachowski E, Reade MC, Cooper DJ, Bellomo R (2011). "Dynamic lactate indices as predictors of outcome in critically ill patients." *Critical Care*, **15**(5), R242.

Ooms J (2014). "The jsonlite Package: A Practical and Consistent Mapping Between JSON Data and R Objects." *arXiv preprint arXiv:1403.2805.*

Pollard TJ, Johnson AE, Raffa JD, Celi LA, Mark RG, Badawi O (2018). "The eICU Collaborative Research Database, a freely available multi-center database for critical care research." *Scientific data*, **5**, 180178.

Singer M, Deutschman CS, Seymour CW, Shankar-Hari M, Annane D, Bauer M, Bellomo R, Bernard GR, Chiche JD, Coopersmith CM, Hotchkiss RS, Levy MM, Marshall JC, Martin GS, Opal SM, Rubenfeld GD, van der Poll T, Vincent JL, Angus DC (2016). "The Third International Consensus Definitions for Sepsis and Septic Shock (Sepsis-3)." *JAMA*, **315**(8), 801–810.

Subbe C, Kruger M, Rutherford P, Gemmel L (2001). "Validation of a modified Early Warning Score in medical admissions." *QJM: An International Journal of Medicine*, **94**(10), 521–526.

Therneau TM, Grambsch PM (2000). *Modeling survival data: extending the Cox model.* Springer, New York.

Therneau TM, Lumley T (2015). "Package 'survival'." *R Top Doc*, **128**, 112.

Thoral PJ, Peppink JM, Driessen RH, Sijbrands EJG, Kompanje EJO, Kaplan L, Bailey H, Kesecioglu J, Cecconi M, Churpek M, Clermont G, van der Schaar M, Ercole A, Girbes ARJ, Elbers PWG, Force obotAUMCDAC, the SCCM/ESICM Joint Data Science Task (2021). "Sharing ICU Patient Data Responsibly Under the Society of Critical Care Medicine/European Society of Intensive Care Medicine Joint Data Science Collaboration: The Amsterdam University Medical Centers Database (AmsterdamUMCdb) Example." *Critical Care Medicine*, **Latest Articles**.

Van Beest PA, Brander L, Jansen SP, Rommes JH, Kuiper MA, Spronk PE (2013). "Cumulative lactate and hospital mortality in ICU patients." *Annals of intensive care*, **3**(1), 6.

Villar J, Pérez-Méndez L, Blanco J, Añón JM, Blanch L, Belda J, Santos-Bouza A, Fernández RL, Kacmarek RM, Spanish Initiative for Epidemiology and Therapies for ARDS (SIESTA) Network S (2013). "A universal definition of ARDS: the PaO2/FiO2 ratio under a standard ventilatory setting—a prospective, multicenter validation study." *Intensive Care Medicine*, **39**(4), 583–592.

Vincent JL, Moreno R, Takala J, Willatts S, De Mendonça A, Bruining H, Reinhart C, Suter P, Thijs LG (1996). "The SOFA (Sepsis-related Organ Failure Assessment) score to describe organ dysfunction/failure."

Wang RZ, Sun CH, Schroeder PH, Ameko MK, Moore CC, Barnes LE (2018). "Predictive Models of Sepsis in Adult ICU Patients." In *2018 IEEE International Conference on Healthcare Informatics (ICHI)*, pp. 390–391.

Wang S, McDermott MB, Chauhan G, Ghassemi M, Hughes MC, Naumann T (2020). "Mimic-extract: A data extraction, preprocessing, and representation pipeline for mimic-iii." In *Proceedings of the ACM Conference on Health, Inference, and Learning*, pp. 222–235.

Wickham H, Hester J (2020). *readr: Read Rectangular Text Data.* R package version 1.4.0, URL https://CRAN.R-project.org/package=readr.

Affiliation:

Nicolas Bennett[1]
ETH Zürich
Seminar for Statistics
Rämistrasse 101
CH-8092 Zurich
E-mail: nicolas.bennett@stat.math.ethz.ch

Drago Plečko[1]
ETH Zürich
Seminar for Statistics
Rämistrasse 101
CH-8092 Zürich
E-mail: drago.plecko@stat.math.ethz.ch

Ida-Fong Ukor
Monash Health
Department of Anaesthesia and Perioperative Medicine
246 Clayton Road
Clayton VIC 3168
E-mail: dr.ukor@gmail.com

Nicolai Meinshausen
ETH Zürich
Seminar for Statistics
Rämistrasse 101
CH-8092 Zürich
E-mail: meinshausen@stat.math.ethz.ch

---

[1]These authors contributed equally.

Peter Bühlmann
ETH Zürich
Seminar for Statistics
Rämistrasse 101
CH-8092 Zürich
E-mail: peter.buehlmann@stat.math.ethz.ch